


Advanced OpenGL

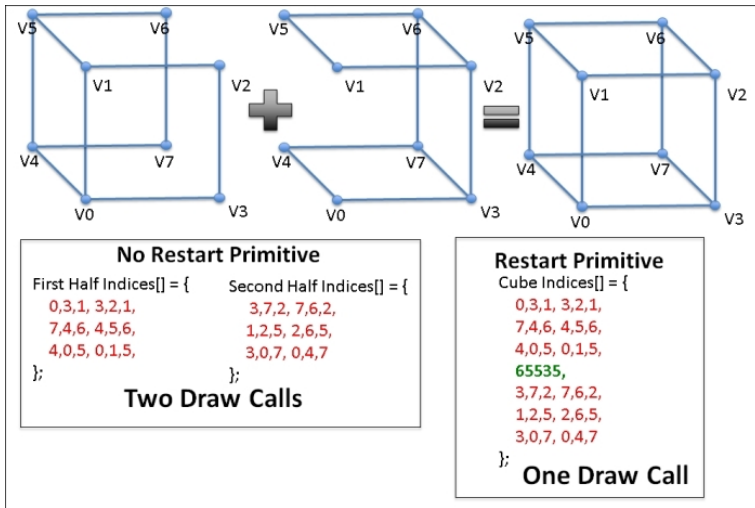
Tomáš Milet

Brno University of Technology, Faculty of Information Technology
Božetěchova 1/2. 602 00 Brno - Královo Pole
imilet@fit.vutbr.cz



22. října 2021





https://www.packtpub.com/mapt/book/application_development/9781849695527/3/ch03lv11sec36/

- Primitive restart index can be use to terminated triangle strip.
- If we want to draw a lot of triangle strips that are not connected, we can use PRI.
- Without PRI, we would have to use `glDrawElements` for each triangle strip.
- or we would have to degenerated some primitives.
- PRI is used as element of element array buffer.
- Primitive assembly is terminated after reaching PRI.

```
glEnable(GL_PRIMITIVE_RESTART);  
glPrimitiveRestartIndex(0xffffffff);
```

- There is a lot of variants of following commands:

```
glDrawArrays glDrawElements
```

- Instancing
- Indirect Draw
- Multi Draw

```
glDrawArrays, glMultiDrawArrays  
  
glDrawArraysInstanced, glDrawArraysInstancedBaseInstance,  
  
glDrawArraysIndirect, glMultiDrawArraysIndirect  
  
glDrawElements, glDrawRangeElements, glMultiDrawElements  
  
glDrawElementsBaseVertex, glDrawRangeElementsBaseVertex, glMultiDrawElementsBaseVertex  
glDrawElementsInstanced, glDrawElementsInstancedBaseInstance, glDrawElementsInstancedBaseVertex,  
glDrawElementsInstancedBaseVertexBaseInstance  
  
glDrawElementsIndirect, glMultiDrawElementsIndirect
```



http://epicbattles.wikia.com/wiki/File:UEBS_Chunk.jpg


- Instancing draws same mesh in multiple instances, each with different transformation matrix.
- Instance id can be access in shader through built-in variable: `gl_InstanceID`
- Instance id can be used as index into array of transformation matrices or materials.

```
#version 460

layout(location=0)in vec4 Position;
uniform mat4 MVP[100];

void main() {
    gl_Position=MVP[gl_InstanceID]*Position;
}
```

```
glBindVertexArray(VAO);
glDrawArraysInstanced(GL_TRIANGLES,0,NumVertices,NumInstances);
glBindVertexArray(0);
```



- Draw call can be stored in buffers.
- There is no need for CPU synchronization.
- It can be combined with compute shader, atomic counters, transform feedback, ...

```
//init
glGenBuffers(1, &IndirectBuffer);
glBindBuffer(GL_DRAW_INDIRECT_BUFFER, IndirectBuffer);
unsigned Data[4] = {100, 1, 0, 0}; //command
glBufferData(GL_DRAW_INDIRECT_BUFFER, sizeof(unsigned) * 4,
             Data, GL_DYNAMIC_DRAW);

//...
//fill buffer from GPU
//...

//draw
glBindBuffer(GL_DRAW_INDIRECT_BUFFER, IndirectBuffer);
glDrawArraysIndirect(GL_TRIANGLES, NULL);
```

- Multi draw call fuses a lot of single draw calls into one.
- All draw calls parameters can be stored in buffers and generated on GPU.
- Frustum Culling on GPU - number of instances of visible object $\neq 0$

```
//init
glGenBuffers(1, &IndirectBuffer);
glBindBuffer(GL_DRAW_INDIRECT_BUFFER, IndirectBuffer);
glBufferData(GL_DRAW_INDIRECT_BUFFER, sizeof(unsigned) * 5 * NumSpheres,
             NULL, GL_DYNAMIC_COPY);
//...
glDispatchCompute(NumSpheres / WorkGroupSize.x + 1, 1, 1);
//...
//draw
glBindVertexArray(VAO);
glBindBuffer(GL_DRAW_INDIRECT_BUFFER, IndirectBuffer);
glMultiDrawElementsIndirect(GL_TRIANGLES, GL_UNSIGNED_INT, NULL,
                             NumSpheres, sizeof(unsigned) * 5);
glBindVertexArray(0);
```

- We often need to draw a lot of objects with different textures.
- This leads to texture switching in between draw commands and to greater number of draw commands.
- Texture binding is not cheap and there is limited number of texture units.
- Texture atlas has its own problems (mipmapping, color bleeding, ...)
- There is an extension: [ARB_bindless_texture](#) that solves a lot of problems
- This extension allows shader to switch textures directly.
- Textures are not explicitly bind to any texture unit.

```
#define MAX_TEXTURES 512
GLuint textures[MAX_TEXTURES]; //list of textures
glCreateTextures(GL_TEXTURE_2D, MAX_TEXTURES, textures); //create textures
for (unsigned i=0; i<MAX_TEXTURES; ++i) { //loop over textures
    glTextureStorage2D(textures[i], 1, GL_RGBA32F, 1024, 1024); //texture allocation
    glTextureSubImage2D(textures[i], 0, 0, 0, 1024, 1024,
        GL_RGBA, GL_UNSIGNED_BYTE, data[i]); //data upload
    ...
}
```

```
GLuint64 handles[MAX_TEXTURES]; //list of handles to textures
for (unsigned i=0; i<MAX_TEXTURES; ++i) {
    handles[i] = glGetTextureHandleARB(textures[i]);
    glMakeTextureHandleResidentARB(handles[i]);
}
```

```
glProgramUniformHandleui64vARB(
    programId, //program id
    uniformLocation, //location of uniform variable
    MAX_TEXTURES, //number of handles
    handles); //handles
```

```
#version 440 core

#extension GL_ARB_bindless_texture : require

#define MAX_TEXTURES 512

layout(bindless_sampler)uniform sampler2D textures[MAX_TEXTURES];
flat out sampler2D sampler;

struct Material{
    uint   textureId;
    vec3   color;
};

layout(std430,binding=0)buffer MaterialArray{Material materials[]};

layout(location=3) in uint materialID;

void main(){
    sampler = textures[materials[materialID]];
    gl_Position=...;
}
```

```
#version 440 core


#extension GL_ARB_bindless_texture : require

layout(location=0) out vec4 fColor;

flat in sampler2D sampler;

in vec2 vTexCoord;

void main() {
    fColor=texture(sampler,vTexCoord);
}
```



- Since OpenGL 4.5, Khronos added large extension into core profile - Direct State Access.
- DSA adds a lot of new functions for direct manipulation of objects.
- Bind to modify paradigm will be deprecated in the future.

Old approach:

- 1 Backup object id that is currently bound to target.
- 2 Bind new object to target.
- 3 Execute operation on target.
- 4 Rebind backup to target.

New approach:

- 1 Execute operation on object.

- DSA is easier and follow common sense.
- DSA does not break bindings.
- DSA is faster, less state switching.
- Target is replaced with object id in most OpenGL functions.
- DSA exists as extension for a long time - it is supported on all GPUs

Bind to modify:

```
GLint old;
glGenRenderbuffers(1, &RBO);
glGetIntegerv(GL_RENDERBUFFER_BINDING, &old);
glBindRenderbuffer(GL_RENDERBUFFER, RBO);
glRenderbufferStorage(GL_RENDERBUFFER, internalFormat, width, height);
glBindRenderbuffer(GL_RENDERBUFFER, old);
```

Direct state access:

```
glCreateRenderbuffers(1, &RBO);
glNamedRenderbufferStorage(RBO, internalFormat, width, height);
```

Bind to modify:

```
glBindFramebuffer(GL_READ_FRAMEBUFFER, fbo);  
glBindFramebuffer(GL_DRAW_FRAMEBUFFER, 0);  
glBlitFramebuffer(0, 0, width, height, 0, 0, width, height,  
    GL_COLOR_BUFFER_BIT, GL_NEAREST);
```

Direct state access:

```
glBlitNamedFramebuffer(fbo, 0, 0, 0, width, height, 0, 0, width, height,  
    GL_COLOR_BUFFER_BIT, GL_NEAREST);
```

- Some OpenGL action have 2 or even 3 alternative commands.
- A lot of functions is redundant.
- Confusion in texture, vao commands

```
glFramebufferTexture  
glFramebufferTexture1D;  
glFramebufferTexture2D;  
glFramebufferTexture3D;  
glFramebufferTextureLayer;  
glNamedFramebufferTexture;  
glNamedFramebufferTextureLayer;
```

```
glVertexArrayBindingDivisor  
glVertexAttribDivisor  
glVertexBindingDivisor
```

```
glTexImage2D  
glTexStorage2D  
glTextureStorage2D
```

- Graphics application contains a lot of different effects (bump mapping, shadows, parallax mapping, ...)
- Different effects require different shaders.
- Functionality switching can be performed on different levels.
- Switching program path using uniform variable.
- Subroutines switching
- Shader program pipelines switching
- Shader program switching

The fastest functionality switch can be performed directly inside shader. However, there is per-invocation overhead. For example, vertex shader has to evaluate condition for every vertex (thread divergence).

```
#version 430
uniform uint method;
void main() {
    switch(method) {
        case 1:
            normalMapping(...);
            break;
        case 2:
            paralaxMapping(...);
            break;
        case 3:
            ...
    }
}
```

It is the second fastest functionality switch. It is function pointer switching.

#version 440

```
layout(location=0)out vec4 fColor;

subroutine vec4 getColorSubroutine();
subroutine vec4 rotateColorSubroutine(vec4);

subroutine (getColorSubroutine) vec4 redColor(){return vec4(1,0,0,1);}
subroutine (getColorSubroutine) vec4 greenColor(){return vec4(0,1,0,1);}

subroutine (rotateColorSubroutine) vec4 rotate1Left (vec4 c){return c.yzwx;}
subroutine (rotateColorSubroutine) vec4 rotate2Left (vec4 c){return c.zwxy;}
subroutine (rotateColorSubroutine) vec4 rotate3Left (vec4 c){return c.wxyz;}
subroutine (rotateColorSubroutine) vec4 reverse (vec4 c){return c.wzyx;}

subroutine uniform getColorSubroutine getColor;//fce pointer
subroutine uniform rotateColorSubroutine rotateColor[3];//array of fce pointers

uniform uint rotateIndex=0;

void main(){
    fColor=rotateColor[rotateIndex](getColor());
}
```

```
//all subroutines
GLuint s10=glGetSubroutineIndex(program, GL_FRAGMENT_SHADER, "redColor");
GLuint s11=glGetSubroutineIndex(program, GL_FRAGMENT_SHADER, "greenColor");
GLuint s12=glGetSubroutineIndex(program, GL_FRAGMENT_SHADER, "rotate1Left");
GLuint s13=glGetSubroutineIndex(program, GL_FRAGMENT_SHADER, "rotate2Left");
GLuint s14=glGetSubroutineIndex(program, GL_FRAGMENT_SHADER, "rotate3Left");
GLuint s15=glGetSubroutineIndex(program, GL_FRAGMENT_SHADER, "reverse");

//uniform offsets for subroutines
GLuint sul0=glGetSubroutineUniformLocation(program, GL_FRAGMENT_SHADER, "getColor");
GLuint sull=glGetSubroutineUniformLocation(program, GL_FRAGMENT_SHADER, "rotateColor");

glUseProgram(program);
//list of all slots for subroutines
GLuint sl[4];
sl[sul0+0]=s11;
sl[sull+0]=s15;
sl[sull+1]=s12;
sl[sull+2]=s13;
glUniformSubroutinesuiv(GL_FRAGMENT_SHADER, 4, sl);
```

- Program pipelines can be used for shader stage switching (vertex,fragment,geometry,...).
- Inputs and outputs of stages has to match, they will not be checked.
- Pipeline switching is cheaper than program switching.
- It is useful in situation where there is a lot of similar shader programs.

```
GLuint vs,gs,fs;
//This function creates shader program that contains only one shader stage.
vs=glCreateShaderProgramv(GL_VERTEX_SHADER,1,&vstext);
fs=glCreateShaderProgramv(GL_FRAGMENT_SHADER,1,&fstext);
//...

GLuint pipeline;
glGenPipelines(1,&pipeline);
glBindPipelines(pipeline);

//...
glUseProgramStages(pipeline, GL_VERTEX_SHADER_BIT, vs);
glUseProgramStages(pipeline, GL_GEOMETRY_SHADER_BIT, program); //select geometry shader from shader program
glUseProgramStages(pipeline, GL_FRAGMENT_SHADER_BIT, fs);
```


- Geometry shader is located after vertex shader (after tessellation).
- Geometry shader is executed per primitive - it has access to all vertices of a primitive.
- It can generate new geometry or modify current geometry.
- It can transform point into quad (usefull for particle simulation).
- It can be used for several effects (shadow volumes, particle systems, ...).
- Geometry Instancing.
- Transform feedback.

- Type of inputs and outputs has to be specified inside geometry shader.

```
layout(points,invocations=N) in; //input primitive will be point.  
//points, lines, lines_adjacency, triangles, triangles_adjacency  
//geometry shader will be executed N times on each point.
```

- Type of output primitive and number of output vertices has to be specified.

```
layout(triangle_strip,max_vertices=4) out; //output primitive and max number of vertices  
//points, line_strip, triangle_strip
```



```
#version 430
```


```
layout(points) in;
```

```
layout(triangle_strip,max_vertices=4) out;
```


```
void main() {  
    gl_Position=mvp*(gl_in[0].gl_Position+vec4(-1,-1,0,0));  
    EmitVertex();  
    gl_Position=mvp*(gl_in[0].gl_Position+vec4(-1,+1,0,0));  
    EmitVertex();  
    gl_Position=mvp*(gl_in[0].gl_Position+vec4(+1,-1,0,0));  
    EmitVertex();  
    gl_Position=mvp*(gl_in[0].gl_Position+vec4(+1,+1,0,0));  
    EmitVertex();  
    EndPrimitive();  
}
```

```
#version 430
layout(points) in;
layout(triangle_strip,max_vertices=4) out;
void main() {
    gl_Position=vec4(-1,-1,0,1);EmitVertex();
    gl_Position=vec4(-1,+1,0,1);EmitVertex();
    gl_Position=vec4(+1,-1,0,1);EmitVertex();
    gl_Position=vec4(+1,+1,0,1);EmitVertex();
    EndPrimitive();
}
```

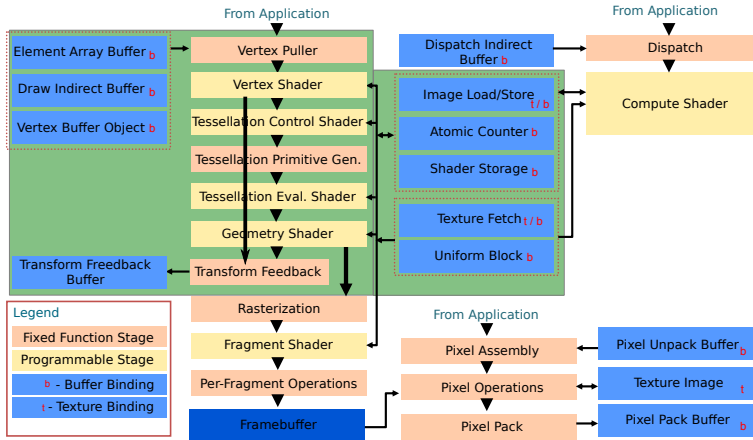
```
glGenVertexArrays(1,&emptyVAO);
//...
glBindVertexArray(emptyVAO); //activate VAO
glDrawArrays(GL_POINTS,0,1);
glBindVertexArray(0); //deactivate VAO
```

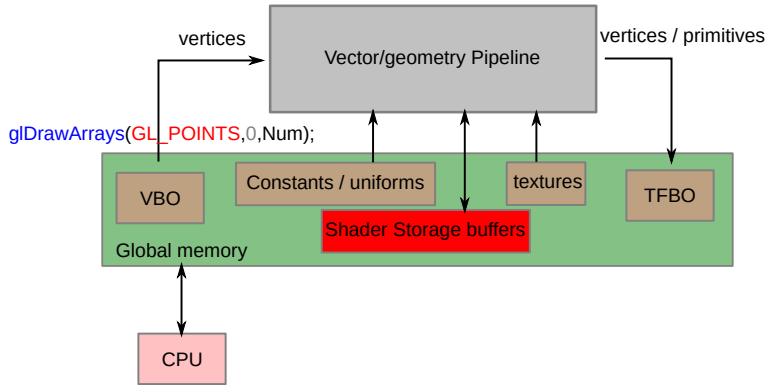


```
#version 330
layout(triangles) in;
layout(triangle_strip,max_vertices=10) out;
uniform mat4 MVP,M; //matrices
uniform vec4 LightPosition; //light position
void main() {
    vec4 LP=M*LightPosition;
    vec4 p[6];
    p[0]=gl_in[0].gl_Position; //triangle vertices
    p[1]=gl_in[1].gl_Position;
    p[2]=gl_in[2].gl_Position;
    p[3]=vec4(gl_in[0].gl_Position.xyz*LP.w-LP.xyz,0); //triangle points projected to infinity
    p[4]=vec4(gl_in[1].gl_Position.xyz*LP.w-LP.xyz,0);
    p[5]=vec4(gl_in[2].gl_Position.xyz*LP.w-LP.xyz,0);
    vec3 N=normalize(cross((p[1]-p[0]).xyz,(p[2]-p[0]).xyz));
    float Distance=dot(N,LP.xyz)-dot(N,p[0].xyz);
    if(Distance<=0){ //flip volume inside out
        vec4 c=p[0];p[0]=p[1];p[1]=c;
        c=p[3];p[3]=p[4];p[4]=c;
    }
    gl_Position=MVP*p[0];EmitVertex();
    gl_Position=MVP*p[1];EmitVertex();
    gl_Position=MVP*p[3];EmitVertex();
    gl_Position=MVP*p[4];EmitVertex();
    gl_Position=MVP*p[5];EmitVertex();
    gl_Position=MVP*p[1];EmitVertex();
    gl_Position=MVP*p[2];EmitVertex();
    gl_Position=MVP*p[0];EmitVertex();
    gl_Position=MVP*p[5];EmitVertex();
    gl_Position=MVP*p[3];EmitVertex();
    EndPrimitive();
}
```



- Transform feedback stops rendering pipeline before rasterisation and sends primitives into buffers.
- It can be used in vertex shader and geometry shader.
- It can be combined with rendering (streams).





```
const char*Vayrings[]={ "Out1", "Out2"};
glTransformFeedbackVaryings(Program,2,Varyings,GL_SEPARATE_ATTRIBS);
glLinkProgram(Program);

//...

glBindBufferBase(GL_TRANSFORM_FEEDBACK_BUFFER,0,Buffer1);
glBindBufferBase(GL_TRANSFORM_FEEDBACK_BUFFER,1,Buffer2);

glEnable(GL_RASTERIZER_DISCARD);//nebudeme rasterizovat
//...
glBeginTransformFeedback(GL_TRIANGLES);
glDrawArrays(...);
glEndTransformFeedback();
```

We have to link shader program with marked varyings. c++:

```
//list of variables that will be written into buffer(s).
const char*ResetVaryings[]={ "vPosition", "vVelocity", "vMass"};
//set varyings and interleaving
glTransformFeedbackVaryings(ResetProgram,3,ResetVaryings, GL_INTERLEAVED_ATTRIBS);
//relink shader program
glLinkProgram(ResetProgram);
```


glsl:

```
#version 330


layout(location=0)out vec2 vPosition;//particle position
layout(location=1)out vec2 vVelocity;//particle velocity
layout(location=2)out float vMass;//particle mass
//...
void main() {
    vPosition = vec2(0); //init position
    vVelocity = vec2(cos(VelAngle), sin(VelAngle)) * VelSize; //init velocity
    vMass = Noise(MassSeed + uint(gl_VertexID), MinMass, MaxMass); //init mass
}
```



- Tessellation splits one primitive into more joint sub primitives.
- It can be used to refine details of a geometry.
- Tessellation is located between vertex shader and geometry shader.
- It is composed of three parts:
 - Control Shader
 - Primitive generation/tessellation
 - Evaluation Shader
- There is one new primitive type - **GL_PATCHES**

```
glPatchParameteri(GL_PATCH_VERTICES,10); //set number of vertices of patch to 10  
glDrawArrays(GL_PATCHES,0,200); //draw 20 patches (each is composed of 10 vertices)
```



- CS controls level of tessellation.
- It computes control vertices.
- CS is invoked as many times as there is output vertices in output patch primitive.
- A invocation number is stored in built-in variable `gl_InvocationID`.
- We can synchronise threads in CS using `barrier()`.



**Quads****Triangles**

- `layout ({isolines,triangles,quads}) in;`
- `gl_TessLevelOuter(4),gl_TessLevelInner(2)`

```
#version 430
```

```
// number of output vertices in output patch primitive ==  
// number of threads per patch primitive
```

```
layout(vertices=3) out;
```

```
uniform vec2 TessLevelInner; //there are two inner and
```

```
uniform vec4 TessLevelOuter; //four outer levels of tessellation
```

```
void main() {
```

```
    //number of elements in gl_in depends on GL_PATCH_VERTICES
```

```
    //number of elements in gl_out depends on layout(vertices=n)out;
```

```
    //In this case, number of elements in gl_in and gl_out is the same.
```

```
    gl_out[gl_InvocationID].gl_Position=gl_in[gl_InvocationID].gl_Position; //copy
```

```
    if(gl_InvocationID==0){ //the first thread sets tessellation levels
```

```
        gl_TessLevelOuter[0]=TessLevelOuter[0];
```

```
        gl_TessLevelOuter[1]=TessLevelOuter[1];
```

```
        gl_TessLevelOuter[2]=TessLevelOuter[2];
```

```
        gl_TessLevelOuter[3]=TessLevelOuter[3];
```


```
        gl_TessLevelInner[0]=TessLevelInner[0];
```

```
        gl_TessLevelInner[1]=TessLevelInner[1];
```

```
    }
```

```
}
```

- It sets type of tessellated primitive: **isolines**, **triangles**, **quads**
- It computes coordinates of tessellated vertices.
- **gl_TessCoord** variable holds barycentric, uv, or normalized coords of tessellated vertices inside of primitive.
- Evaluation shader is executed for every tessellated vertex.
- Tessellated primitives continue their path to geometry shader.



- Tessellated quad
- Computation of coordinates of tessellated vertices

```
#version 430
```

```
layout (quads) in;
```

```
void main() {
```

```
    vec4 A=mix(gl_in[0].gl_Position,gl_in[1].gl_Position,gl_TessCoord.x);
```

```
    vec4 B=mix(gl_in[3].gl_Position,gl_in[2].gl_Position,gl_TessCoord.x);
```

```
    gl_Position=mix(A,B,gl_TessCoord.y);
```

```
}
```

```
// Vertex shader
#version 430
void main() {
    gl_Position = mvp*position;
}

// Control shader
#version 430
layout(vertices=16) out;

void main() {
    gl_out[gl_InvocationID].gl_Position =
    gl_in[gl_InvocationID].gl_Position;
    if(gl_InvocationID == 0) {
        gl_TessLevelInner[0] = gl_TessLevelInner[1] =
        gl_TessLevelOuter[0] = gl_TessLevelOuter[1] =
        gl_TessLevelOuter[2] = gl_TessLevelOuter[3] = 64;
    }
}
```

```
// Evaluation shader
#version 430
layout(quads, ccw) in;

vec4 bernstein(float t) {
    return vec4((1-t)*(1-t)*(1-t), 3*t*(1-t)*(1-t), 3*t*t*(1-t), t*t*t);
}

void main() {
    vec4 bu = bernstein(gl_TessCoord.x);
    vec4 bv = bernstein(gl_TessCoord.y);
    vec4 position = vec4(0, 0, 0, 0);
    for(int y = 0; y < 4; ++y){
        for(int x = 0; x < 4; ++x){
            position += bu[x]*bv[y]*gl_in[4*y + x].gl_Position;
        }
    }
    gl_Position = position;
}
```

```
#version 430

//output vertex shader attributes
out vec4 vAttrib;
gl_Position



//input control shader attributes
in vec4 vAttrib[]; //attribute from vertex shader, size == GL_PATCH_VERTICES
gl_in[].gl_Position; //a position attribute from vertex shader

//output control shader attributes
out vec4 cAttrib[]; //size == layout(vertices=size)out
patch out mat4 cM; //once per output patch primitive

//input evaluation shader attributes
in vec4 cAttrib[]; //size == layout(vertices=size)out
patch in mat4 cM; //once per input patch

//output evaluation shader attributes
out vec3 eNormal;

//input geometry shader attributes
in vec3 eNormal[]; //size == 2 for line, size == 3 for triangle, ...
```



$$K = \begin{pmatrix} 1 & 0 & 0 & S_x \\ 0 & 1 & 0 & S_y \\ 0 & 0 & 1 & S_z \\ 0 & 0 & 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} r & 0 & 0 & 0 \\ 0 & r & 0 & 0 \\ 0 & 0 & r & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} u_x & v_x & w_x & 0 \\ u_y & v_y & w_y & 0 \\ u_z & v_z & w_z & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

Control Shader

```
#version 400
layout(vertices=1)out;
patch out mat4 K;
void main() {
    gl_TessLevelOuter[0]=1;
    gl_TessLevelOuter[1]=64;
    gl_TessLevelOuter[2]=1;
    gl_TessLevelOuter[3]=1;
    gl_TessLevelInner[0]=1;
    gl_TessLevelInner[1]=1;
    vec4 TT[3];
    TT[0]=gl_in[0].gl_Position;
    TT[1]=gl_in[1].gl_Position;
    TT[2]=gl_in[2].gl_Position;
    float t01=length((TT[0]-TT[1]).xyz);
    float t02=length((TT[0]-TT[2]).xyz);
    float t12=length((TT[1]-TT[2]).xyz);
    float s=t01+t02+t12;
    float r=sqrt((s/2-t01)*(s/2-t02)*(s/2-t12)*s/2)*2/s;
    t01/=s;
    t02/=s;
    t12/=s;
    vec3 C=TT[0].xyz*t12+TT[1].xyz*t02+TT[2].xyz*t01;
    vec3 x=normalize(TT[0].xyz-C);
    vec3 y=normalize(TT[1].xyz-C);
    vec3 z=normalize(cross(x,y));
    y=normalize(cross(z,x));
    K=mat4(vec4(x,0)*r,vec4(y,0)*r,vec4(z,0)*r,vec4(C,1));
}
```

Evaluation Shader

```
#version 400

#define MY_PI 3.14159265359

layout(isolines)in;

uniform mat4 V;
uniform mat4 P;

patch in mat4 K;

void main() {
    float Angle=gl_TessCoord.x*MY_PI*2;
    vec4 PP=vec4(cos(Angle),sin(Angle),0,1);
    gl_Position=P*V*K*PP;
}
```


Outer level:

- Outer levels of edges of neighbor faces should be the same to prevent T-joints.
- Transform control points to the screen (screen-space distances).
- Divide edge length by maximal edge length.

Inner level:

- Average/maximum of outer levels
- ...

- Reads/Writes from/to images inside of shader stage.
- There is new data type - image.
- An image is one layer of texture (one layer of mipmap, ...).
- There are atomic operation for images.
- Atomic operation are supported only for certain internal formats (integer, one channel).
- Image units (equal 8 units) - similar to texture units (equal 80 units).
- Store operations are side effect. They disable early fragment tests in fragment shader - it can be reenabled.



This fragment shader emulates stencil buffer:

```
#version 430

layout(location=0)out vec4 fColor;
layout(early_fragment_tests)in; //enable early fragment tests
layout(r32i,location=0)uniform iimage2D myStencil; //integer image
ivec2 Coord=ivec2(gl_FragCoord.xy); //coordinates

void main() {
    if(gl_FrontFacing) //front face
        imageAtomicAdd(myStencil, Coord, +1);
    else
        imageAtomicAdd(myStencil, Coord, -1);
}
```

```
//integer texture
glGenTextures(1, &Image);
glBindTexture(GL_TEXTURE_2D, Image);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, GL_NEAREST);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, GL_NEAREST);
glTexImage2D(GL_TEXTURE_2D, 0, GL_R32I, Widht, Height, 0, GL_RED_INTEGER,
             GL_UNSIGNED_BYTE, NULL);

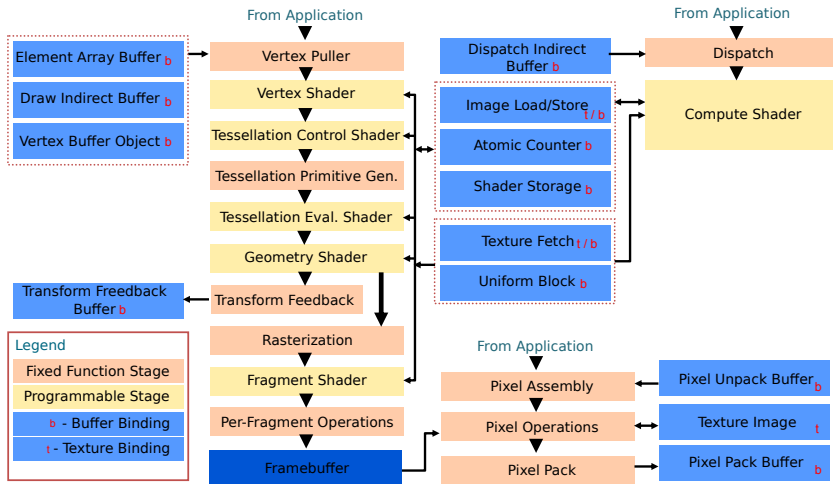
//bind one layer of the texture onto zeroth image unit
glBindImageTexture(0, Image, 0, GL_FALSE, 0, GL_READ_WRITE, GL_R32I);
```

Atomic counter:

- Atomic increment/decrement.
- It can be used as index into shader storage buffers.
- Indirect Draw

Shader Storage Buffer

- Random access memory for shaders
- Atomic operation
- Binding points



```
#version 430

in vec4 vPos;

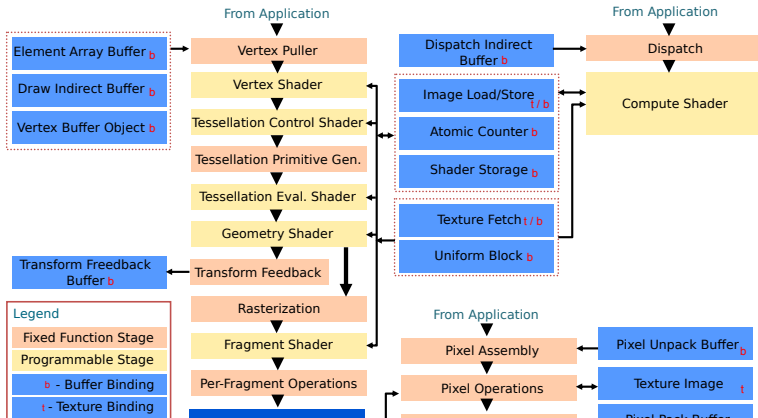
layout(binding=1,offset=0)uniform atomic_uint counter;
layout(std430,binding=0)buffer Output{vec4 data[]};

layout(location=0)vec4 fColor;
//...
void main() {
    int W=atomicCounterIncrement(counter);//increment counter
    data[W*2+0]=ComputeColor(...);//compute color
    data[W*2+1]=vPos;//write position
    //...
}
```


```
unsigned data[4]={0,0,0,0};
GLuint ACB;//identifier of a. counter
glCreateBuffers(1,&ACB);//create a. counter buffer
glNamedBufferData(ACB,sizeof(uint32_t)*4,data,GL_DYNAMIC_DRAW);
glBindBufferBase(GL_ATOMIC_COUNTER_BUFFER,1,ACB);//binding point 1

GLuint SSBO;//identifier of shader storage buffer
glCreateBuffers(1,&SSBO);//reserve identifier
glNamedBufferData(GL_SHADER_STORAGE_BUFFER,sizeof(float)*4*2*Max,
    NULL,GL_DYNAMIC_DRAW);
glClearNamedBufferData(SSBO,GL_R32F,GL_RED,GL_FLOAT,NULL);
glBindBufferRange(GL_SHADER_STORAGE_BUFFER,0,SSBO,0,
    sizeof(float)*4*2*Max);//binding point 0
```

- Compute shader is located outside of common rendering pipeline.
- Compute shader can be used for general computing - GPGPU (other shader are in some way restricted).
- Compute shader share syntax with the rest of shader and can be used on different platforms.



OpenGL Compute Programming Model and Compute Memory Hierarchy



```
#version 430

layout(binding=0)buffer Input{vec4 i[]};; //input buffer
layout(binding=1)buffer Output{vec4 o[]};; //output buffer
layout(local_size_x=32,local_size_y=1,local_size_z=1)in; //work-group size

uniform uint num; //number of elements in buffers

void main() {
    //gl_GlobalInvocationID = gl_WorkGroupID*gl_WorkGroupSize + gl_LocalInvocationID
    uint gid = gl_GlobalInvocationID.x;
    if(gid >= num) return;
    vec3 v = i[gid].xyz;
    o[gid] = vec4(normalize(v), length(v));
}
```

```
uint32_t num=100;
GLuint CompBufferInput;
glCreateBuffers(1,&CompBufferInput);
glNamedBufferData(CompBufferInput,sizeof(float)*4*num,nullptr,GL_STATIC_DRAW);

GLuint CompBufferOutput;
glCreateBuffers(1,&CompBufferOutput);
glNamedBufferData(CompBufferOutput,sizeof(float)*4*num,nullptr,GL_DYNAMIC_COPY);

glBindBufferBase(GL_SHADER_STORAGE_BUFFER,0,CompBufferInput);
glBindBufferBase(GL_SHADER_STORAGE_BUFFER,1,CompBufferOutput);

glUseProgram(ComputeShader);
glUniform1ui(glGetUniformLocation(ComputeShader,"num"),num);

uint32_t workGroupSize = 32;

glDispatchCompute(divRoundUp(num,workGroupSize),1,1);
```

- Synchronization needs to be done between two shaders that operates on same data.
- Synchronization is performed using `glMemoryBarrier()`; command.

```
//This compute shader creates buffer with vertices.
```

```
glDispatchCompute(  
    this->TileCount [ (this->NumLevels-2)*2+0],  
    this->TileCount [ (this->NumLevels-2)*2+1],  
    1);
```

```
//Wait for modification to SSBO
```

```
glMemoryBarrier(GL_SHADER_STORAGE_BARRIER_BIT);
```

```
//Draw generated vertices
```

```
glDrawArrays(...)
```

- Threads in work-group are execute in warp/wavefronts.
- We can synchronize threads that are part of work-group.
- We cannot synchronize threads that are part of different work-groups.
- Threads that are part of warp do not have to be synchronized.
- Synchronization is performed by barrier() command.
- A barrier command has to lie on program path that can be access by all threads in work-group.
- A barrier command stops every thread until all other threads in work-group reach it.

Example:

```
layout(std430, binding=0) readonly buffer SFData{float data[]};  
shared float array[size];  
void main() {  
    //cooperative reading from global memory into local/shared memory.  
    array[gl_LocalInvocationID.x]=data[gl_GlobalInvocationID.x];  
    //wait for all threads  
    barrier();  
    //...  
}
```

- Queries can be used to obtain some information about drawing.
- Number of rasterized samples.
- Number of generated primitives.
- Time of execution of commands.
- Conditional rendering - occlusion culling

```
void glGenQueries(GLsizei n, GLuint * ids);
```

```
void glBeginQuery(GLenum target, GLuint id);
```

```
void glEndQuery(GLenum target, GLuint id);
```

```
void glGetQueryiv(GLenum target, GLuint id, GLint * params);
```

```
//init
GLuint QueryTime;//query
glGenQueries(1,&QueryTime);//nagenerujeme si ID query
GLuint QueryTimePassed=0;//cas v nanosekundach

//start query
glBeginQuery(GL_TIME_ELAPSED,QueryTime);

//commands
//glBindBuffer(...);
//glUseProgram(...);
//glDrawArrays(...);
//...

//end query
glEndQuery(GL_TIME_ELAPSED);//vypneme query

//get data from asynchronous query
glGetQueryObjectuiv(QueryTime,GL_QUERY_RESULT_NO_WAIT,&QueryTimePassed);
```

```
//init
GLuint QuerySample; //query
glGenQueries(1, &QueryTime); //nagenerujeme si ID query

//draw bounding boxes
//start query
glBeginQuery(GL_ANY_SAMPLES_PASSED, QuerySample);
//draw bounding boxes
//glBindBuffer(...);
//glUseProgram(...);
//glDrawArrays(...);
//...
glEndQuery(GL_ANY_SAMPLES_PASSED); //vypneme query

//conditional rendering
//start conditional rendering
glBeginConditionalRender(QuerySample, GL_QUERY_NO_WAIT);
//draw meshes with full details
//glBindBuffer(...);
//glUseProgram(...);
//glDrawArrays(...);
//...
glEndConditionalRender(); //end conditional rendering
```


- Synchronization between OpenGL commands.
- Host synchronization, client synchronization.

```
GLsync sync; //synchronization object
glDispatchCompute(...);
sync=glFenceSync(GL_SYNC_GPU_COMMANDS_COMPLETE, 0);
glClientWaitSync(sync, 0, GL_TIMEOUT_IGNORED); //CPU waits
//...
glDeleteSync(sync);
```

```
GLsync sync; //synchronization object
glDispatchCompute(...);
sync=glFenceSync(GL_SYNC_GPU_COMMANDS_COMPLETE, 0);
glWaitSync(sync, 0, GL_TIMEOUT_IGNORED); //GPU waits
//...
glDeleteSync(sync);
```

```
glFlush();
glFinish();
```

- OpenGL exposes commands for debugging.
- User can define custom messages, callbacks, ...
- Old way: `glGetError`
- OpenGL debugging is supported in debug OpenGL context

```
void glDebugMessageCallback(DEBUGPROC callback, void * userParam);
```

```
void glDebugMessageControl(GLenum, GLenum, GLenum, GLsizei, const GLuint*, ...);
```

```
void glDebugMessageInsert(GLenum, GLenum, GLuint, GLenum, GLsizei, const char*);
```

```
//create window using SDL2
SDL_Init(SDL_INIT_VIDEO|SDL_INIT_TIMER); //init. video
SDL_GL_SetAttribute(SDL_GL_CONTEXT_MAJOR_VERSION, 4);
SDL_GL_SetAttribute(SDL_GL_CONTEXT_MINOR_VERSION, 3);
SDL_GL_SetAttribute(SDL_GL_CONTEXT_PROFILE_MASK,
    SDL_GL_CONTEXT_PROFILE_CORE);
//set debug context
SDL_GL_SetAttribute(SDL_GL_CONTEXT_FLAGS, SDL_GL_CONTEXT_DEBUG_FLAG);
SDL_CreateWindow("OpenGL s debug kontekstem", SDL_WINDOWPOS_CENTERED,
    SDL_WINDOWPOS_CENTERED, 1024, 768,
    SDL_WINDOW_OPENGL|SDL_WINDOW_SHOWN|SDL_WINDOW_FULLSCREEN);

//custom callback
void MyDebug(GLenum Source, GLenum Type, GLuint Id, GLenum severity,
    GLsizei Length, const GLchar* Message, void* UserParam) {
    std::cerr<<"MyDebug: "<<Message<<std::endl;
}

glEnable(GL_DEBUG_OUTPUT); //enable debugging
glDebugMessageCallback(MyDebug, NULL); //set callback
```

- <http://www.opengl.org/sdk/docs/>
- <http://www.opengl.org/documentation/glsl/>
- <https://wiki.libsdl.org/FrontPage>

Thank you for your attention! Questions?